

AI Engineering & Consultation — Selected Work

Three projects demonstrating end-to-end AI system design: from RAG pipelines and LLM agent infrastructure to applied data engineering. Built with Python, LangGraph, MCP, FastAPI, and Claude.

Python · AI/ML · LLM Agents August 2026

Project 1

RAG Chatbot — Document-Grounded Q&A System

A retrieval-augmented generation chatbot that lets users upload documents and ask questions, receiving answers grounded in those documents with inline source citations.

TECHNOLOGY STACK

LangGraph

Claude API

ChromaDB


HuggingFace Embeddings

Streamlit

LangChain

Python asyncio

ARCHITECTURE WORKFLOW



Two-stage RAG pipeline: offline document ingestion with local embeddings, and LangGraph-orchestrated query-time retrieval + generation.

WHAT I BUILT

- Designed and implemented a complete RAG pipeline with LangGraph's **StateGraph** for orchestrating retrieve → generate flow
- Used local HuggingFace sentence-transformers for embeddings — indexing and retrieval run fully offline with zero API cost
- Persistent ChromaDB vector store survives restarts; supports PDF, TXT, and Markdown document ingestion
- Streaming answer generation via Claude with inline source citations from retrieved passages
- Built conversation memory (last 6 turns) for contextual follow-up questions
- Provider-agnostic architecture — embeddings can be swapped to Voyage AI or OpenAI with a single function change


KEY METRICS

2-stage
LangGraph pipeline

Top-4
Chunks per retrieval

\$0
Embedding cost (local)

3 formats
PDF / TXT / MD



Project 2

MCP MySQL Server — AI Agent Database Infrastructure

A Model Context Protocol server that gives AI agents (Claude Desktop, Claude Code) structured, safe access to MySQL databases — with a companion eval suite for deterministic validation.

TECHNOLOGY STACK

MCP SDK

Claude Agents

PyMySQL


Python asyncio

SQLAlchemy

pytest

dotenv

ARCHITECTURE WORKFLOW



MCP server exposes 6 tools with safety-first defaults; the eval suite drives the server end-to-end via a real MCP client session.

WHAT I BUILT

- Built a production MCP server exposing 6 tools for database interaction — reads enabled by default, writes and DDL explicitly opt-in via environment flags
- Implemented SQL statement classification that rejects mismatched operations (read_query won't accept INSERT, write_query won't accept DDL)
- Built a companion deterministic eval suite that launches the server as a subprocess, drives it through real MCP client sessions, and validates across 9 test categories
- Designed composable scorer functions (`expect_json` , `expect_error` , `values_contain`) for declarative test case definitions
- Published as both a Claude Desktop skill (zip upload) and a Claude Code plugin (GitHub-linked, auto-updating)

KEY METRICS

6 tools

MCP-exposed
operations

23/23

Eval pass rate

9 categories

Test coverage areas

Read-only

Default safety mode

Project 3

Uber Eats Tracker — Personal Spend Analytics Platform

A data engineering project that imports Uber Eats order data from multiple sources (CSV, JSON, receipt emails), stores it in a normalized database, and surfaces spending analytics through a web dashboard and CLI.

TECHNOLOGY STACK

FastAPI

SQLAlchemy 2.0

Chart.js


Typer CLI

PyMySQL

IMAP / email parsing

pytest

ARCHITECTURE WORKFLOW



Multi-source ETL pipeline: three importers normalize data into a cents-precision ORM; analytics engine serves insights via dashboard, API, and CLI.

WHAT I BUILT

- Built three pluggable importers (CSV, JSON, .eml email) with a shared normalization layer and cents-precision money storage to avoid float drift
- Designed the email importer with regex-based receipt parsing tolerant of formatting changes, with clear error boundaries
- Implemented a Gmail auto-sync script using IMAP that fetches new Uber Eats receipts since the last known order date and inserts deduplicated records
- Built 5 analytics functions: summary stats, monthly breakdown, top restaurants, top items, day-of-week spending patterns
- Delivered through both a FastAPI web dashboard (with Chart.js visualizations) and a Typer CLI for scripting and automation

- Full test suite with in-memory SQLite for fast, isolated tests

KEY METRICS

3 importers
CSV / JSON / Email

7 endpoints
REST API

5 analytics
Spend insights

Auto-sync
Gmail IMAP → MySQL

About this portfolio. These projects are open-source and available on GitHub. Each was designed, built, and tested end-to-end — from architecture decisions through implementation and deployment. Code samples and full repositories available upon request.

Core competencies demonstrated: RAG pipeline design, LLM agent infrastructure (MCP), vector databases, data engineering & ETL, API design, evaluation frameworks, full-stack development, Python ecosystem fluency.